

LAMP-Merge: Long-tail-Aware Medical Prototype Merging for One-shot Asynchronous Clinical Models

Anonymous submission

Abstract

Multicenter medical models are typically trained asynchronously across hospitals with heterogeneous acquisition devices, patient populations, and diagnostic-label coverage. We study single-shot asynchronous medical model merging, where each client uploads its locally trained model only once and the server constructs a global model without accessing raw medical images or conducting federated optimization rounds. We identify aggregation-induced prediction collapse as the central failure mode of general-purpose model-merging methods. Furthermore, we show that long-tailed medical priors can make majority-class collapse deceptively competitive in terms of raw accuracy, thereby masking this failure. To address this issue, we propose LAMP-Merge, a long-tail-aware medical prototype-merging framework that reconstructs a diagnostic prototype head from client-side class statistics and calibrates long-tail prevalence using a bounded prior bias. Across five medical datasets, four vision-model families, three client-population sizes, and three levels of non-IID skew, LAMP-Merge matches or surpasses the strongest non-LAMP baseline in 57 of 60 client-averaged settings, while improving balanced accuracy, macro-F1, and prediction-collapse diagnostic metrics.

Index Terms: medical model merging, asynchronous fusion, predictive collapse, long-tail calibration, privacy-preserving learning

I Introduction

Medical artificial intelligence increasingly requires multicenter collaboration because no single hospital can cover the full disease spectrum, rare cases, and imaging conditions (Sheller et al. 2020; Rieke et al. 2020). Yet privacy requirements, device heterogeneity, and institutional workflows make centralizing raw images impractical (Kaissis et al. 2020a). A realistic alternative is for hospitals to train locally and share limited post-training information.

This setting requires single-shot asynchronous model merging. Hospitals may complete training at different times and cannot necessarily participate in repeated global updates. The server must therefore construct one deployable diagnostic model from exactly one upload per client, without accessing raw medical images.

Medical imaging makes the problem especially difficult because hospitals differ not only in device domain but also in diagnostic-class availability (Guan and Liu 2022; Zech




Figure 1: In multicenter healthcare settings, individual clients typically observe only a subset of the global diagnostic space. Conventional parameter-level model-merging methods can therefore cause the merged global model’s predictions to concentrate on a small number of dominant classes.

et al. 2018). Local models are partial experts: a client may provide strong evidence for common diagnoses while lacking rare ones. Moreover, medical labels are often long-tailed, so a useful merge must preserve distributed class-specific evidence while respecting genuine prevalence (Buda, Maki, and Mazurowski 2018).

We propose LAMP-Merge, a long-tail-aware medical prototype-merging framework. It recasts global parameter averaging as class-level diagnostic-evidence reconstruction: each client uploads class prototypes and counts computed with a shared reference extractor, and the server constructs a global prototype classifier with bounded prevalence calibration. The procedure requires one upload and never exposes raw images, per-sample features, logits, or predictions.

Our main contributions are as follows:

- We formalize single-shot asynchronous medical image model merging, where independently trained clients provide one upload for global model construction.
- We identify heterogeneous diagnostic-class availability and long-tailed disease prevalence as the defining structures of this setting.




Figure 2: Local models from multicenter hospitals typically contain only partial diagnostic knowledge, while raw medical images cannot be centrally shared. LAMP-Merge reconstructs a privacy-friendly global medical classifier through diagnostic prototype reconstruction and long-tail prevalence calibration, without accessing raw data or requiring multi-round federated optimization.

- We propose LAMP-Merge, which reconstructs a global diagnostic classifier from client-side class prototypes and prevalence statistics to mitigate prediction collapse without exposing raw images.

II Related Works

Model Merging

Model-merging methods commonly integrate checkpoints in parameter, task-vector, or geometric spaces. Weight averaging and Model Soups interpolate model parameters (Wortsman et al. 2022); Task Arithmetic, TIES-Merging, and DARE operate on task deltas and sign conflicts (Ilharco et al. 2023; Yadav et al. 2023; Yu et al. 2024); Fisher merging, RegMean, Model Stock, Breadcrumbs, Iso-C, FreeMerge, and RobustMerge further introduce curvature, mean-statistic, subspace, or robustness assumptions (Matena and Raffel 2022; Jin et al. 2025; Jang, Yun, and Han 2025; Davari and Belilovsky 2024; Marczak et al. 2025; Zheng and Wang 2025; Zeng et al. 2025).

However, these variables are still global parameters, task vectors, or update directions rather than class-level diagnostic evidence. **They do not ask whether a client has reliable supervision for a specific medical class.** Under incomplete local class coverage and long-tailed prevalence, reliable directions from clients that observed a class can be averaged with absent or conflicting directions from clients that did not, leading to aggregation-induced prediction collapse.

Medical Model Merging

Medical fusion studies often exploit domain structures such as irregular clinical time series, physiological sensor graphs,

and image-EHR fusion. mTAN, Raindrop, and MedFuse show that medical tasks benefit from modeling clinical sampling mechanisms, modality heterogeneity, and patient-level temporal relations (Shukla and Marlin 2021; Zhang et al. 2022; Hayat, Geras, and Shamout 2023).

However, these methods usually fuse patient-level inputs or clinical representations rather than independently trained image-classifier checkpoints, and they often require patient samples, timestamps, multimodal records, electronic health records, or server-side validation feedback. MedMerge is a relevant training-free medical knowledge-transfer framework (Wei et al. 2026); however, it targets cross-modal adaptation for medical vision-language models rather than asynchronous checkpoint merging for the same classification task. **Thus, existing medical fusion routes are mismatched with our privacy and single-upload protocol constraints.** LAMP-Merge keeps the medical structure uploadable by using only class-level prototypes and prevalence counts computed locally by clients.

Federated Learning

Federated and decentralized learning perform collaborative training through communication protocols. FedAvg and its variants rely on repeated server-client updates (McMahan et al. 2017; Kaissis et al. 2020b); Swarm Learning and Gossip Learning reduce centralization but still require participants to continue training, exchanging models, and optimizing from intermediate states (Wan et al. 2021; Hegedüs, Danner, and Jelasity 2019).

However, our problem is single-shot asynchronous model merging after local training has finished: hospitals may not remain online, receive global checkpoints repeatedly, or run

additional optimization rounds. **The limitation is therefore communication feasibility rather than only privacy.** LAMP-Merge requires one upload of a local model interface and class-level statistics, enabling the server to construct a unified diagnostic model without a multi-round feedback loop.

III Method

Motivation and Overall Idea

Under partial local class coverage, parameter-level merging can turn several non-degenerate local predictors into a globally collapsed predictor. Because medical prevalence is often long-tailed, such collapse can still appear competitive in raw accuracy. A useful merger must therefore recover rare-class discrimination without discarding genuine prevalence.

LAMP-Merge addresses this tension with two components. Diagnostic prototype reconstruction aggregates only class-supported evidence in a shared reference space to recover global multiclass directions. Long-tail prevalence calibration then applies a bounded prior bias estimated from client-side counts, preserving clinically meaningful majority prevalence without replacing those directions. The resulting one-shot protocol requires neither additional federated training nor raw-data exposure.

Overview

The core framework of LAMP-Merge is illustrated in figure2. Its central idea is to recast medical model merging from global parameter averaging into the reconstruction of class-conditional diagnostic evidence, supplemented by class-prior information.

Notably, LAMP-Merge neither requires the server to access clients' raw images nor relies on a server-side validation set, per-sample features, or per-sample predictions. Clients upload only class-level statistics; thus, the framework preserves a single-shot asynchronous communication boundary, distinguishing it from federated learning procedures that require repeated model broadcasts and local retraining.

The framework comprises two key components: a class-level diagnostic prototype reconstruction module and a long-tail prevalence calibration module. The prototype reconstruction module restores multiclass discriminative directions across the complete global diagnostic space, after which the calibration module retains appropriate class-prior information.

Diagnostic Prototype Reconstruction Module

Assume K medical clients. Client i owns a private labeled dataset $D_i = \{(x, y)\}$ drawn from a class-skewed medical distribution over C diagnosis classes. Let $D_{i,c}$ denote the subset of client i belonging to diagnosis class c . In medical partitions, $D_{i,c}$ can be empty for many classes. Thus, a client can be reliable for a subset of diagnoses but unreliable for unseen classes. A class-agnostic parameter merge has no explicit mechanism to preserve this local reliability structure.

Therefore, LAMP-Merge first reconstructs a class-wise diagnostic head in the shared reference feature space. Let

ϕ_0 be the frozen reference backbone instantiated from the shared model configuration before local training, and let $T(\cdot)$ be the evaluation transform. For each diagnosis class c , client i computes a reference-space class prototype

$$\mu_{i,c} = \frac{1}{n_{i,c}} \sum_{(x,y) \in D_i} \mathbf{1}[y = c] \phi_0(T(x)), \quad (1)$$

where $n_{i,c} = |D_{i,c}|$ is the class support count. If $n_{i,c} = 0$, the client uploads zero support for that class and contributes no prototype evidence.

The server converts support counts into evidence weights:

$$e_{i,c} = (n_{i,c} + 1)^\gamma \mathbf{1}[n_{i,c} > 0], \quad (2)$$

where $\gamma = 0.45$ is the default evidence exponent. The normalized reliability of client i for class c is

$$\alpha_{i,c} = \frac{e_{i,c}}{\sum_{j=1}^K e_{j,c}}. \quad (3)$$

The global diagnosis prototype is then

$$p_c = \sum_{i=1}^K \alpha_{i,c} \mu_{i,c}. \quad (4)$$

Finally, LAMP-Merge synthesizes a cosine prototype classifier:

$$w_c = s \frac{p_c}{\|p_c\|_2}, \quad (5)$$

where $s = 20$ is the default head scale. This module addresses a key limitation of existing methods: it preserves the classes for which each client has supporting evidence and reconstructs a global multiclass diagnostic head.

Long-tail Prevalence Calibration

After obtaining this newly reconstructed diagnostic head, each client additionally uploads class-prevalence counts $m_{i,c}$ to account for the long-tailed prevalence inherent in real-world medical tasks. The server estimates the global class prior

$$\pi_c = \frac{\sum_{i=1}^K m_{i,c}}{\sum_{k=1}^C \sum_{i=1}^K m_{i,k}}. \quad (6)$$

We measure dominant-class imbalance by

$$r = C \max_c \pi_c. \quad (7)$$

The module is activated only when $r > \tau$, where $\tau = 2.5$ by default. When activated, it adds a centered log-prior bias

$$b_c = \lambda \left(\log \pi_c - \frac{1}{C} \sum_{k=1}^C \log \pi_k \right), \quad (8)$$

with default maximum calibration strength $\lambda = 5.0$. The final score is

$$\text{score}_c(x) = w_c^\top \phi_0(T(x)) + b_c. \quad (9)$$

The bias is centered, bounded, and attached to the prototype classifier; it calibrates prevalence without replacing the multi-class diagnostic directions reconstructed by the Diagnostic Prototype Reconstruction Module.




Figure 3: Representative training images from the five medical image datasets used in the experiments. The class IDs are shown only to illustrate category diversity within each dataset.




Figure 4: Dataset-level client-average accuracy in the module ablation. Each dataset contains 12 cells from four vision-model families and three client counts; each cell is averaged over three Dirichlet skew levels.

Experiments

Experimental Settings

1) Evaluation Datasets: We evaluate the merged models on multiple representative medical image benchmarks. These datasets cover key medical imaging modalities, including blood-cell microscopy, dermoscopy, abdominal CT organ slices, and ultrasound images, and involve cell morphology recognition, skin-lesion diagnosis, organ-structure recognition, and ultrasound-structure discrimination. The datasets used in the experiments include:

- **Blood:** an eight-class blood-cell microscopy dataset for hematological morphology classification, focusing on fine-grained cellular appearance and inter-class morphological similarity.
- **Derma:** a seven-class dermoscopic skin-lesion dataset, characterized by strong long-tailed prevalence and visually subtle lesion categories.
- **Organ-C:** an 11-class abdominal CT dataset in the coronal view, evaluating organ-category recognition under grayscale cross-sectional imaging.
- **Organ-S:** an 11-class abdominal CT dataset in the sagittal view, complementing Organ-C with a different anatomical imaging plane.
- **Ultrasound:** an eight-class ultrasound image dataset, covering structure recognition under modality-specific speckle noise and low-contrast boundaries.

These benchmark datasets cover diverse imaging settings in medical image classification, providing a systematic testbed for evaluating the cross-modality robustness,

class-discriminative capability, and task generalization of merged models. We conduct systematic experiments on these datasets to analyze the performance of LAMP-Merge across different medical image types and class-distribution conditions.

2) Baselines: We compare LAMP-Merge with twelve model-merging baselines covering the main technical routes of current model merging: avg(Wortsman et al. 2022), ties(Yadav et al. 2023), dare_linear(Yu et al. 2024), dare_ties(Yu et al. 2024), regmean(Jin et al. 2025), fisher(Matena and Raffel 2022), breadcrumbs(Davari and Belilovsky 2024), model_stock(Jang, Yun, and Han 2025), from(Ilharco et al. 2023), iso_c(Marczak et al. 2025), free_merge(Zheng and Wang 2025), and robustmerge(Zeng et al. 2025). These baselines include weight averaging, task-vector pruning, sign-conflict resolution, curvature or mean-statistic correction, stock-style model interpolation, isotropic subspace merging, Fourier-domain fusion, and direction-robust merging.

3) Merging Settings: Experiments follow a single-shot asynchronous model-merging protocol. Each client first independently trains a local model on its own medical imaging data, then computes the class-level statistics required by LAMP-Merge locally, including class-support counts, reference-space class prototypes, and class-prevalence counts. These statistics are uploaded once together with an interface to the local checkpoint. The server subsequently performs a single model-merging step using only the uploaded checkpoints and class-level statistics, and evaluates the resulting merged model on a global test set.

We evaluate four vision-model families: ResNet, ConvNeXt, ViT-Tiny, and Swin-Tiny. For each dataset-model pair, we construct non-IID client partitions with $K \in \{3, 5, 7\}$ clients and Dirichlet skew parameter $\beta \in \{0, 0.01, 0.1\}$, using seed 42 (Hsu, Qi, and Brown 2019). A client-average result fixes the dataset, model, and K , and then averages over the three β values.

4) Implementation Environment: All experiments were run on a server with two Intel(R) Xeon(R) Silver 4210 CPUs, four NVIDIA GeForce RTX 2080 Ti GPUs, and the conda environment MM with Python 3.10.20, PyTorch 2.6.0+cu118, and CUDA runtime 11.8.

Comparison With the State of the Arts

The main comparison uses client-average ACC: for each fixed dataset, model, and client number, we first average over the three β skew levels and then compare merging methods. A cell is counted as successful when LAMP-Merge is greater than or equal to the strongest non-LAMP baseline in that cell.

Tables 1–4 report the full client-average results. Across the 60 client-average cells, LAMP-Merge wins or ties the strongest non-LAMP baseline in 57 cells, corresponding to 95.0%. This result indicates that the gain is not an accidental effect of one client number, one partition skew, or one model family; instead, it appears consistently across multiple medical image modalities and backbone architectures.

Ablation Study

Module-level ablations use exactly the same experimental setup as the main experiments. We compare three configurations: full LAMP-Merge, diagnostic prototype reconstruction only, and weight averaging augmented only with long-tail prevalence calibration. Fig. 4 shows that both modules provide substantial performance improvements.

1) Effectiveness of Diagnostic Prototype Reconstruction

To further identify which information within diagnostic prototype reconstruction is necessary, we conduct controlled substitutions across the full setting. All configurations retain long-tail prevalence calibration and vary only the prototype semantics or support-count weighting used in diagnostic prototype reconstruction.

- **LAMP-Merge** uses $p_c = \sum_i \alpha_{i,c} \mu_{i,c}$, where $e_{i,c} = (n_{i,c} + 1)^\gamma \mathbf{1}[n_{i,c} > 0]$ determines the class-support reliability.
- **Classifier-head aggregation** replaces the reference-space class prototype by the local classifier-head direction, i.e., $\mu_{i,c} \leftarrow h_{i,c}$.
- **Global-feature mean** removes class conditioning by setting $\mu_{i,c} \leftarrow g$ for all c , where g is the class-agnostic global reference-feature mean.
- **Support-only synthetic head** removes image evidence by replacing each prototype with a fixed random unit direction u_c , i.e., $p_c \leftarrow u_c, \|u_c\|_2 = 1$.
- **Shuffled-label prototype** preserves prototype magnitudes but permutes diagnostic semantics, i.e., $p_c \leftarrow$

$p_{\sigma(c)}$, where σ is a random permutation over diagnosis labels.

- **Uniform client weight** replaces support-weighted aggregation with equal weights among clients containing class c , i.e., $\alpha_{i,c} \leftarrow \mathbf{1}[n_{i,c} > 0] / \sum_j \mathbf{1}[n_{j,c} > 0]$.
- **Binary support only** replaces fine-grained support and prevalence counts with class-presence indicators, i.e., $n_{i,c}, m_{i,c} \leftarrow \mathbf{1}[n_{i,c} > 0]$.
- **Global client-size weight** uses the total client support $N_i = \sum_k n_{i,k}$ rather than class-specific support, i.e., $e_{i,c} \leftarrow (N_i + 1)^\gamma \mathbf{1}[n_{i,c} > 0]$.

Fig. 4 and Table 5 show that neither an arbitrary head nor a class-agnostic statistic explains the gain. Replacing reference-space prototypes or class-specific support weights consistently reduces ACC, confirming that diagnostic semantics and class-level reliability are jointly necessary.

2) Effectiveness of Long-tail Prevalence Calibration

We further replace the prior term inside long-tail prevalence calibration. This group changes only the class prior used by the calibration bias, and the second column gives the corresponding formula.

- **LAMP-Merge** estimates the empirical class prior as $\pi_c = \sum_i m_{i,c} / \sum_k \sum_i m_{i,k}$.
- **No prevalence calibration** removes the calibration bias by setting $b_c \leftarrow 0$.
- **Uniform prevalence prior** removes long-tail prevalence by setting $\pi_c \leftarrow 1/C$, where C is the number of diagnosis classes.
- **Smoothed prevalence prior** uses $\pi_c \leftarrow \sum_i (m_{i,c} + 1) / \sum_k \sum_i (m_{i,k} + 1)$, preserving the same prior structure with additive smoothing.

Table 6 shows that removing prevalence calibration or replacing the real prior with a uniform prior reduces mean ACC, confirming that long-tailed prevalence is part of the full method. The smoothed prior remains close to the official method, indicating that LAMP-Merge does not depend on accidental fluctuations of exact counts; bounded calibration remains effective as long as the prior preserves the real long-tail structure.

Hyperparameter Analysis Experiment

Fig. 5 evaluates the prototype-head scale s and calibration strength λ over 180 raw cells and 60 client-average cells, spanning five datasets, four vision backbones, three client counts, and three Dirichlet skew levels. The default setting is $s = 20$ and $\lambda = 5$.

Both curves remain stable over broad intervals. Small s is insufficient to scale prototype logits, whereas bounded λ calibrates prevalence without inducing majority-class prediction. Thus, the gain arises from the two components rather than narrow hyperparameter tuning.

New-metric Diagnostic Experiment

Although the main experiments use ACC for direct comparison, long-tailed medical distributions can allow single-class or few-class predictors to achieve high ACC. We therefore introduce additional metrics. Let y_n and $\hat{y}_n$ denote




Figure 5: Full-scope hyperparameter sensitivity. The default setting $s = 20$ and $\lambda = 5$ lies in the stable high-accuracy plateau.

Table 5: Internal ablation of diagnostic prototype reconstruction. All settings cover 180 raw cells and 60 client-average cells.

Setting	Mean Acc	Δ vs. LAMP
LAMP-Merge	0.6209	—
Classifier-head aggregation	0.2579	$\downarrow 0.3630$
Global-feature mean	0.2645	$\downarrow 0.3564$
Support-only synthetic head	0.1137	$\downarrow 0.5073$
Shuffled-label prototype	0.2002	$\downarrow 0.4207$
Uniform client weight	0.5839	$\downarrow 0.0370$
Binary support only	0.5517	$\downarrow 0.0692$
Global client-size weight	0.5944	$\downarrow 0.0266$

Table 6: Internal ablation of long-tail prevalence calibration.

Setting	Mean Acc	Δ vs. LAMP
LAMP-Merge	0.6209	—
No prevalence calibration	0.5879	$\downarrow 0.0330$
Uniform prevalence prior	0.5879	$\downarrow 0.0330$
Smoothed prevalence prior	0.6209	0.0000

the true and predicted labels of the n -th test sample. Let $q_c = N^{-1} \sum_n \mathbf{1}[\hat{y}_n = c]$ be the predicted-class distribution, $p_c = N^{-1} \sum_n \mathbf{1}[y_n = c]$ be the true class distribution, and $\mathcal{C}_{\text{test}} = \{c \mid \sum_n \mathbf{1}[y_n = c] > 0\}$ be the set of classes observed in the test set. We use the following diagnostic metrics.

- $\text{BA} = |\mathcal{C}_{\text{test}}|^{-1} \sum_{c \in \mathcal{C}_{\text{test}}} \text{Recall}_c$. This is balanced accuracy, or macro recall, and measures whether all observed diagnosis classes are recalled.
- $\text{MacroF1} = |\mathcal{C}_{\text{test}}|^{-1} \sum_{c \in \mathcal{C}_{\text{test}}} \text{F1}_c$. This measures class-balanced precision-recall quality and penalizes methods that only preserve majority-class predictions.
- $\rho = \max_c q_c$. This is the prediction-collapse ratio and measures how strongly predictions concentrate on a single diagnosis class.
- $C_{\text{eff}} = \exp(-\sum_c q_c \log q_c)$. This is the effective number of predicted classes and measures how many diagnosis classes are actually used by the model.
- $\text{TV}(q, p) = \frac{1}{2} \sum_c |q_c - p_c|$. This is the total variation

Table 7: Prediction-collapse diagnostics averaged over five medical datasets and four backbones per dataset.

Metric	LAMP-Merge	Best reference	Δ vs. LAMP
Accuracy $\uparrow$	0.5885	0.3156	$\uparrow +0.2729$
BA / macro recall $\uparrow$	0.5747	0.1810	$\uparrow +0.3937$
Macro F1 $\uparrow$	0.5352	0.1123	$\uparrow +0.4229$
Collapse ratio $\downarrow$	0.2294	0.7477	$\downarrow -0.5184$
Effective classes $\uparrow$	8.0390	2.1360	$\uparrow +5.9030$
Pred-true TV $\downarrow$	0.1510	0.6168	$\downarrow -0.4658$

distance between predicted and true class distributions and measures whether the output distribution matches the test-set prevalence.

Table 7 compares LAMP-Merge with the best-performing baseline for each metric. LAMP-Merge consistently and substantially outperforms the existing methods. This indicates that its accuracy gains do not arise from more severe single-class prediction collapse, but rather from more complete global diagnostic discrimination.

Visualization Experiment

To visualize the prediction-distribution changes measured by the new diagnostic metrics, we provide a representative predicted-class distribution in Fig. 6.

t-SNE Visualization Experiment

We use t-SNE only for analysis, not for server-side selection or fusion. Fig. 7 shows a joint comparison on Blood with ResNet, $K = 3$, and $\beta = 0.01$. Each point is the output-probability representation of a test image, and formal merging methods are projected into the same two-dimensional space to compare the class structure of their predictive representations.

In-depth Analysis Experiment

1) analysis of the geometry Finally, we conduct an in-depth analysis of the geometry induced by diagnostic prototype reconstruction. This experiment corresponds




Figure 6: Predicted class distributions on Ultrasound, averaged over the four visual backbones, $K \in \{3, 5, 7\}$, and the three Dirichlet skew levels. The generic merge concentrates predictions on one class, while LAMP-Merge recovers a broader class allocation aligned with the test distribution.

to the prototype-replacement settings in the internal ablation and explains why reference-space class prototypes and class-level support counts jointly provide stable gains. Let $d(a, b) = 1 - \cos(a, b)$ denote cosine distance, and let $\mathcal{I}_c = \{i \mid n_{i,c} > 0\}$ denote the set of clients that contain diagnosis class c . The geometric quantities are defined as follows.

- $D_{\text{pair}} = \frac{2}{C(C-1)} \sum_{c < d} d(p_c, p_d)$. This is the mean distance between global prototypes of different diagnosis classes, measuring whether the reconstructed head forms separable class-level diagnostic directions.
- $D_{\text{nn}} = \frac{1}{C} \sum_c \min_{d \neq c} d(p_c, p_d)$. This is the mean nearest-class distance, measuring the most vulnerable decision margin around each diagnosis prototype.
- $A_{\text{client}} = \frac{1}{C} \sum_c \text{Avg}_{i < j, i, j \in \mathcal{I}_c} \cos(\mu_{i,c}, \mu_{j,c})$. This is the within-class consistency among client prototypes, measuring whether clients that observe the same diagnosis produce compatible evidence in the shared reference space.
- $A_{\text{proto}} = \frac{1}{C} \sum_c \sum_{i \in \mathcal{I}_c} \alpha_{i,c} \cos(p_c, \mu_{i,c})$. This is the alignment between the server-side global prototype and the contributing client prototypes, measuring whether aggregation preserves client-side class evidence.

Fig. 8 reports these four quantities for the controlled prototype substitutions in the internal ablation. The full method jointly preserves nonzero inter-class separation and high cross-client consistency and prototype-client alignment, rather than maximizing a single geometric value. Together with the accuracy drops in Table 5, this joint profile shows that the benefit arises from class-semantic reference-space prototypes weighted by class-specific support.

The individual geometry values should therefore be interpreted as mechanism diagnostics, not standalone objectives: extreme separation or perfect alignment can be produced by class-agnostic or mismatched directions. LAMP-Merge is effective because its four quantities are simultaneously consistent with class semantics and the observed full-scope accuracy gains.

2) Additional Theoretical Analysis Let μ_c^* denote the true class mean in the shared reference space. If client class features have bounded second moments, the prototype estimation error satisfies

$$\mathbb{E} \|p_c - \mu_c^*\|_2^2 \leq \sum_i \alpha_{i,c}^2 \frac{\sigma_c^2}{n_{i,c}} + \text{Bias}_c^2. \quad (10)$$

This upper bound shows that support counts directly control the variance of class-prototype estimation. A client that does not observe class c should not contribute evidence for that class, and a client with more samples should contribute more, but only sublinearly.

Let the true margin between classes c and d for sample x be

$$\Delta_{c,d}(x) = (\mu_c^*)^\top z - (\mu_d^*)^\top z. \quad (11)$$

When the prototype error satisfies

$$\|p_c - \mu_c^*\|_2 + \|p_d - \mu_d^*\|_2 < \frac{\Delta_{c,d}(x)}{\|z\|_2}, \quad (12)$$

the estimated prototypes preserve the class order between c and d .

The role of long-tail prevalence calibration is to introduce long-tail priors in a bounded manner. For prevalence counts $m_{i,c}$, the class prior is

$$\pi_c = \frac{\sum_i m_{i,c}}{\sum_k \sum_i m_{i,k}}, \quad (13)$$

and the centered log-prior bias is

$$b_c = \lambda \left(\log \pi_c - \frac{1}{C} \sum_k \log \pi_k \right). \quad (14)$$

As long as λ is bounded, long-tail prevalence calibration does not replace the discriminative directions reconstructed by diagnostic prototype reconstruction; it only calibrates the output toward the real class prevalence. This explains why using long-tail prevalence calibration alone is insufficient, while full LAMP-Merge remains both non-collapsed and prevalence-aware.

Conclusion

We studied one-shot asynchronous medical model merging, where independently trained hospital models must be fused without federated optimization rounds and without exposing raw medical images to the server. Our empirical analysis shows that generic model-merging methods suffer from aggregation-induced predictive collapse under incomplete local diagnostic coverage. We further show that, on long-tailed medical datasets, such collapse can be masked by raw accuracy. We therefore introduce the Long-tail Prevalence Calibration module and perform secondary evaluation using macro recall, macro-F1, and predictive-distribution diagnostics (Buda, Maki, and Mazurowski 2018; Brodersen et al. 2010).

LAMP-Merge addresses these two properties through a concise two-module design. Future work will extend its applicability in two directions: first, from medical image classification to segmentation, detection, and multimodal diagnostic tasks; and second, toward stronger privacy-preserving




Figure 7: Joint t-SNE visualization of output-probability representations on Blood. From left to right: LAMP-Merge, weight averaging, TIES-Merging, and DARE-Linear.




Figure 8: Prototype-geometry comparison on the full benchmark. The figure summarizes pairwise distance, nearest-class distance, prototype consistency, and prototype-client alignment.

mechanisms, such as differential privacy and secure aggregation, to further reduce potential privacy risks associated with class-level statistics. We will also explore deployment and validation in larger-scale real-world multicenter health-care settings.

References

Brodersen, K. H.; Ong, C. S.; Stephan, K. E.; and Buhmann, J. M. 2010. The Balanced Accuracy and Its Posterior Distribution. In *2010 20th International Conference on Pattern Recognition*, 3121–3124.

Buda, M.; Maki, A.; and Mazurowski, M. A. 2018. A systematic study of the class imbalance problem in convolutional neural networks. *Neural Networks*, 106: 249–259.

Davari, M.; and Belilovsky, E. 2024. Model Breadcrumbs: Scaling Multi-Task Model Merging with Sparse Masks. arXiv:2312.06795.

Guan, H.; and Liu, M. 2022. Domain Adaptation for Medical Image Analysis: A Survey. *IEEE Transactions on Biomedical Engineering*, 69(3): 1173–1185.

Hayat, N.; Geras, K. J.; and Shamout, F. E. 2023. MedFuse: Multi-modal fusion with clinical time-series data and chest X-ray images. arXiv:2207.07027.

Heged\Hus, I.; Danner, G.; and Jelasity, M. 2019. Gossip Learning as a Decentralized Alternative to Federated Learning. In Pereira, J.; and Ricci, L., eds., *Distributed Applications and Interoperable Systems*, 74–90. Cham: Springer International Publishing. ISBN 978-3-030-22496-7.

Hsu, T.-M. H.; Qi, H.; and Brown, M. 2019. Measuring the Effects of Non-Identical Data Distribution for Federated Visual Classification. arXiv:1909.06335.

Ilharco, G.; Ribeiro, M. T.; Wortsman, M.; Gururangan, S.; Schmidt, L.; Hajishirzi, H.; and Farhadi, A. 2023. Editing Models with Task Arithmetic. arXiv:2212.04089.

Jang, D.-H.; Yun, S.; and Han, D. 2025. Model Stock: All we need is just a few fine-tuned models. arXiv:2403.19522.

Jin, X.; Ren, X.; Preotiuc-Pietro, D.; and Cheng, P. 2025. Dataless Knowledge Fusion by Merging Weights of Language Models. arXiv:2212.09849.

Kaissis, G. A.; Makowski, M. R.; Rückert, D.; and Braren, R. F. 2020a. Secure, privacy-preserving and federated machine learning in medical imaging. *Nature Machine Intelligence*, 2(6): 305–311.

Kaissis, G. A.; Makowski, M. R.; Rückert, D.; and Braren, R. F. 2020b. Secure, privacy-preserving and federated machine learning in medical imaging. *Nature Machine Intelligence*, 2(6): 305–311.

Marczak, D.; Magistri, S.; Cygert, S.; Twardowski, B.; Bagdanov, A. D.; and van de Weijer, J. 2025. No Task Left Behind: Isotropic Model Merging with Common and Task-Specific Subspaces. arXiv:2502.04959.

Matena, M.; and Raffel, C. 2022. Merging Models with Fisher-Weighted Averaging. arXiv:2111.09832.

McMahan, B.; Moore, E.; Ramage, D.; Hampson, S.; and Arcas, B. A. y. 2017. Communication-Efficient Learning of Deep Networks from Decentralized Data. In Singh, A.; and Zhu, J., eds., *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, volume 54 of *Proceedings of Machine Learning Research*, 1273–1282. PMLR.

Rieke, N.; Hancox, J.; Li, W.; Milletari, F.; Roth, H. R.; Albarqouni, S.; Bakas, S.; Galtier, M. N.; Landman, B. A.;

Maier-Hein, K.; Ourselin, S.; Sheller, M.; Summers, R. M.; Trask, A.; Xu, D.; Baust, M.; and Cardoso, M. J. 2020. The future of digital health with federated learning. *npj Digital Medicine*, 3(1): 171.

Sheller, M. J.; Edwards, B.; Reina, G. A.; Martin, J.; Pati, S.; Kotrotsou, A.; Milchenko, M.; Xu, W.; Marcus, D.; Colen, R. R.; and Bakas, S. 2020. Federated learning in medicine: Facilitating multi-institutional collaborations without sharing patient data. *Scientific Reports*, 10(1): 10933.

Shukla, S. N.; and Marlin, B. M. 2021. Multi-Time Attention Networks for Irregularly Sampled Time Series. arXiv:2101.10318.

Warnat-Herresthal, S.; Schultze, H.; Shastri; and so on. 2021. Swarm Learning for decentralized and confidential clinical machine learning. *Nature*, 594(7862): 265–270.

Wei, Y.; Zhu, Q.; Xia, Z.; Wang, L.; Zou, Y.; Li, Y.; Luo, Y.; and Du, B. 2026. MedMerge: A Training-Free Model Merging Framework for Medical Knowledge Transfer into Vision–Language Models. *IEEE Transactions on Radiation and Plasma Medical Sciences*, 1–1.

Wortsman, M.; Ilharco, G.; Gadre, S. Y.; Roelofs, R.; Gontijo-Lopes, R.; Morcos, A. S.; Namkoong, H.; Farhadi, A.; Carmon, Y.; Kornblith, S.; and Schmidt, L. 2022. Model soups: averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In Chaudhuri, K.; Jegelka, S.; Song, L.; Szepesvari, C.; Niu, G.; and Sabato, S., eds., *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, 23965–23998. PMLR.

Yadav, P.; Tam, D.; Choshen, L.; Raffel, C.; and Bansal, M. 2023. TIES-Merging: Resolving Interference When Merging Models. arXiv:2306.01708.

Yu, L.; Yu, B.; Yu, H.; Huang, F.; and Li, Y. 2024. Language Models are Super Mario: Absorbing Abilities from Homologous Models as a Free Lunch. arXiv:2311.03099.

Zech, J. R.; Badgeley, M. A.; Liu, M.; Costa, A. B.; Titano, J. J.; and Oermann, E. K. 2018. Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: A cross-sectional study. *PLOS Medicine*, 15(11): 1–17.

Zeng, F.; Guo, H.; Zhu, F.; Shen, L.; and Tang, H. 2025. RobustMerge: Parameter-Efficient Model Merging for MLLMs with Direction Robustness. arXiv:2502.17159.

Zhang, X.; Zeman, M.; Tsiligkaridis, T.; and Zitnik, M. 2022. Graph-Guided Network for Irregularly Sampled Multivariate Time Series. arXiv:2110.05357.

Zheng, S.; and Wang, H. 2025. FREE-Merging: Fourier Transform for Efficient Model Merging. arXiv:2411.16815.